Agent rulebook

Answer these six questions before an AI agent goes live, and again whenever you give it a new job. If you can't answer one, the agent isn't ready.

Agent _____ The job it does _____

Owner _____ Last checked _____

What can it read?

Name the inbox, the folder, the spreadsheet. "Our email" is not an answer; "the support inbox, not the directors' inboxes" is.

What can it change or send?

Drafting a reply and sending one are different permissions. So are reading an invoice and marking it paid.

What can it spend or give away?

Put a number on it: a refund limit, a discount limit, a maximum order value. If the honest answer is nothing, write "nothing".

What must it hand to a person?

Complaints, anything legal, any customer who asks for a human, anything over the limit above.

Who gets told?

When it hands something over or does something unusual, a named person hears about it the same day.

How do you switch it off?

One person, one action, written down. Tested before launch, not discovered during a bad afternoon.

- [ ] Off switch tested on _____ by _____
- [ ] The limits above are checked by the system every time the agent acts, not just written into its instructions.
- [ ] Compared this page with what the agent can actually reach. Anything it can reach that isn't on this page has been removed.

Here's one filled in

Example: an online shop's refund agent

Can read	The orders inbox and order history.
Can send	Refund confirmations only. Any other reply is a draft for the owner.
Can refund	Up to £25, on orders under 30 days old.
Hands to the owner	Anything larger or older.
Who's told	The owner gets a daily summary.
Off switch	Anyone on the team can pause it from the dashboard.

Six lines, one per question, and it covers most of what could go wrong. Yours can be this short. What matters is that every line names something specific: an inbox, an amount, a person.

A limit, or an instruction?

A limit that only exists in a setup screen is a hope. A limit the system checks every time the agent acts is a control.

When you're choosing a tool or a builder, ask which kind you're getting. If the answer is "the agent has been told not to", that's an instruction, not a limit, and instructions can be misread.

Already running an agent?

Fill this page in today, then compare it with what the agent can actually reach. **The gap between the two is your risk.**

About to switch one on?

Fill this page in first. It takes an afternoon and costs nothing, and it will change which jobs you're comfortable handing over.

Not sure where your agent's limits should sit?

Bring the job you're thinking of handing to an agent to a free 30-minute call, and I'll help you decide what it should never be allowed to do. You talk to me, not a salesperson.

Marcus Aragbaye, Founder & AI Lead, AegeanPulse

cal.com/aegeanpulse/ai-strategy-consultation aegeanpulse.com

contact@aegeanpulse.com